

PROCESS LOG

Attempt 1 - ElevenLabs (audio)

Tool/version: ElevenLabs, free tier

Voice used: Roger - Laid-Back, Casual, Resonant

Model: Eleven Multilingual v2

Settings: Speed ~50% (default), Stability ~50% (middle of "more variable" ↔ "more stable"), Similarity ~85% (high), Style Exaggeration 0% (none), Speaker boost ON, Output format MP3 44.1kHz (128kbps)

Prompt/script used: script.txt (unmodified)

Free-tier limits hit: none - script fit comfortably within free character allowance

Output: artifact1_coach_narrative_AI_VOICE_elevenlabs.mp3

Attempt 2 - LOVO / Luvvoice (audio, later used to drive the D-ID video)

Tool/version: LOVO (Luvvoice)

Voice used: Jenny (Female), English (United States)

Settings: Rate 0% (default), Volume 0% (default), Pitch 0% (default)

Prompt/script used: script.txt (unmodified)

Note: This audio was not kept as a standalone artifact - it was uploaded into D-ID to generate the lip-synced video (Attempt 3). Two other free TTS tools were tried first and abandoned before landing on LOVO:

- *TTSMaker*: hit a character limit on the guest tier; signing up for a free account still routed toward premium-only voices for download
- *Murf.ai*: generated audio previewed fine, but download was gated behind a paid upgrade
- Both are worth noting as a finding: several tools market a "free tier" that lets you generate/preview audio but paywalls the actual download.

Attempt 3 - D-ID (video, AI avatar lip-sync)

Tool/version: D-ID, free tier

Avatar used: Amber

Audio source: LOVO/Jenny mp3 (Attempt 2) uploaded as the driving audio track

Settings: Sentiment set to "Professional"

Output: artifact2_coach_narrative_AI_VIDEO_did_lovo.mp4

Failure modes observed (for the assignment's "working vocabulary" objective)

Free-tier "bait": TTSMaker and Murf.ai both allow free generation/preview but gate the actual file download behind a paid plan - a real and repeatable finding about how "free" these tools actually are.

Voice/avatar mismatch: pairing the LOVO (Jenny) voice with the D-ID Amber avatar produced a subtle but noticeable incongruity - the voice didn't feel like a natural match for the avatar's appearance, even though neither element was individually flawed.

Excess avatar motion: the D-ID avatar moved more than a real person would while talking - not just facial animation, but visible shoulder/body movement that reads as distinctly synthetic rather than a recorded human.

Emotional flatness in video: the avatar's delivery lacked ordinary conversational emotion (not exaggeration - just the normal variation a real speaker has), even though lip-sync and blinking were both accurate.

Underlying voice quality gap: the LOVO/Jenny voice driving the video was noticeably more robotic than the ElevenLabs voice - flatter delivery, less natural pacing - which mattered more to overall believability than the video mechanics themselves.

Detection failure: a purpose-built voice deepfake detector (deepfakedetection.io) scored the ElevenLabs audio 95% "Likely Authentic," incorrectly classifying fully synthetic speech as real human recording. See `evaluation.pdf` for full detection write-up.

Total time spent

~1.5–2 hours total, including dead ends: script prep (~10 min), ElevenLabs generation (~5 min), failed attempts on TTSMaker and Murf.ai (~20 min), LOVO generation (~10 min), D-ID video generation (~10–15 min), detection check (~10 min), documentation (~30 min).